

## Homework Five, for Thu 2/18

CSE 250B

Your homework must be typeset, and the PDF file must be uploaded to Gradescope by midnight on the due date.

1. *Voting Perceptrons.* The perceptron algorithm, as presented in class, will not converge when given data that is not linearly separable.

One work-around is to simply halt the algorithm after a fixed number (say,  $T$ ) of passes through the data, and to output the vector  $w$  at that time. But this has a drawback: it might be that an update occurred right before termination, and caused the final  $w$  to point in a poor direction.

A better idea is to combine the various  $w$ 's obtained during the training process. In particular, a vector  $w$  that survived for quite a long time is likely to be good. This is the motivation for the *voted perceptron* algorithm. Given data  $(x^{(1)}, y^{(1)}), \dots, (x^{(n)}, y^{(n)}) \in \mathbb{R}^p \times \{-1, 1\}$ , it operates as follows:

- $\ell = 1, c_\ell = 0$
- $w_1 = 0$
- Repeat  $T$  times:
  - Randomly permute the data points
  - For  $i = 1$  to  $n$ :
    - \* If  $(x^{(i)}, y^{(i)})$  is misclassified by  $w_\ell$ :
      - $w_{\ell+1} = w_\ell + y^{(i)}x^{(i)}$
      - $c_{\ell+1} = 1, \ell = \ell + 1$
    - \* Else:
      - $c_\ell = c_\ell + 1$

At the end of this process, we have a collection of linear separators  $w_1, w_2, \dots, w_\ell$  as well as their “survival times”  $c_1, \dots, c_\ell$ . To classify a new point  $x$ , we take the weighted majority vote:

$$\text{sign} \left( \sum_{j=1}^{\ell} c_j \text{sign}(w_j \cdot x) \right).$$

An alternative rule is the simpler *averaged perceptron*: predict  $\text{sign}(w \cdot x)$ , for

$$w = \sum_{j=1}^{\ell} c_j w_j.$$

- (a) Implement the voted perceptron algorithm and try it out on the small 2-d data set `data1.txt`. (Remember to add a constant-valued feature to the data.) Try  $T = 10$  or thereabouts. Show the final decision boundary. Is it linear?
  - (b) After a few passes through the data, the number of classifiers  $\ell$  might grow inconveniently large. Can you devise a good way to downsample them so that at most  $L$  different values of  $w$  are ever stored? Give pseudocode, and redo the previous problem (part (b)) using your technique. Try a higher value of  $T$  to assess how well your downsampling is working.
  - (c) Give pseudocode for the averaged perceptron that avoids keeping track of all  $w$ 's seen (it should maintain at most two  $w$ 's). Show the decision boundary it achieves on the small data set, with  $T = 10$ .
2. *Kernelized perceptron.* Implement the kernel perceptron algorithm for the quadratic and RBF kernels. Show the boundaries obtained for the small data sets `data1.txt` and `data2.txt`. For the RBF kernel, show boundaries for a few settings of the scale parameter  $\sigma$ .